

Highlights

Decision before Feature: Deciding Whether to Execute a Landscape Descriptor Query in Black-box Optimization

Anonymous Author(s)

- Formulates landscape-descriptor acquisition as a decision made before descriptor computation.
- Uses offline, state-aligned utility labels and algorithm-agnostic search behavior from native optimizer histories.
- Specifies train-only model selection and threshold freezing before any held-out benchmark evaluation.

Decision before Feature: Deciding Whether to Execute a Landscape Descriptor Query in Black-box Optimization

Anonymous Author(s)^a

^aAffiliation withheld for anonymous review,

ARTICLE INFO

Keywords:

black-box optimization
algorithm selection
landscape analysis
trajectory features
resource-aware decision making
function-family generalization

ABSTRACT

Landscape descriptors are widely used to support algorithm selection in black-box optimization, but acquiring them consumes function evaluations, runtime, and memory, while their value may vary across search states. This paper formulates acquisition of a fixed landscape-descriptor query as a decision that precedes feature computation. The proposed Decision-before-Feature protocol reuses algorithm-agnostic, permutation-invariant behavior from complete native optimizer-update histories, constructs state-aligned query-utility labels through paired offline continuations, and freezes a supervised decision rule under nested function-family out-of-fold evaluation. The primary query is a predefined 16-dimensional low-cost descriptor configuration; two larger configurations are reserved for representation-and-cost sensitivity analyses. **TBD** The prevalence and interval for non-beneficial primary-query states remain undetermined. **TBD** The nested-OOF model comparison, frozen held-out BBOB-family evaluation, and B3-versus-T0 increment remain undetermined. **TBD** The end-to-end comparison with the prespecified baselines, including query use, final performance, and resource measures, remains undetermined. **TBD** Performance on BBOB validation, CEC2017, CEC2022, and engineering problems remains undetermined; no transfer claim is made. **TBD** The association between search behavior, search maturity, and query utility remains undetermined and has a descriptive, non-causal interpretation.

1. Introduction

Black-box optimization is used when objective values can be obtained by evaluation but an analytic model, derivative information, or a tractable problem structure is unavailable. In this setting, the choice of optimizer matters: algorithms expose complementary behavior across problem instances, while no member of a practical portfolio is uniformly preferable. Automated algorithm selection addresses this heterogeneity by mapping an instance representation to a portfolio decision. A common workflow evaluates a space-filling design, computes exploratory landscape analysis (ELA) features, and then applies a supervised selector [7, 12]. This workflow makes landscape analysis a prerequisite for selection.

That prerequisite is not innocuous. Landscape characterization comprises families of measures with different assumptions, sampling requirements, search dependence, and computational demands [16, 22]. The objective evaluations used to construct a landscape sample consume budget that could otherwise be assigned to optimization, and the appropriate feature-computation fraction depends on the selection scenario [1]. Contemporary systems therefore count feature evaluations or omit expensive groups under a restricted budget [7]. An ELA-informed portfolio decision may be useful and still fail to justify the resources spent to obtain its descriptors at a particular search state.

At the same time, an optimization prefix already produces information. Search-trajectory networks characterize convergence patterns, recurrent transitions, and attraction regions from ordinary optimizer runs [19]. DynamoRep preserves the longitudinal evolution of population summaries,

while Opt2Vec learns from populations visited during search [3, 13]. Probing trajectories use short objective-value sequences for a subsequent portfolio choice, and later work shows that their evaluation depends on both the classifier and the problem-level split [24, 25]. Reusing evaluated points avoids a separate objective-function sample, but representation and inference still have computational cost. More importantly, problem classification or direct algorithm selection does not establish that a pre-query behavior summary predicts the net value of acquiring an independent descriptor query.

Generalization is a second obstacle. Algorithm-selection performance can appear strong when training and evaluation sets contain closely related instances. Under demanding problem and problem-combination splits, several classical and learned landscape representations did not improve over a single-best-solver baseline in the evaluated affine-BBOB setting [4]. The closest per-run warm-starting study likewise reported that results obtained on BBOB did not transfer directly to YABBOB [14]. Recent reviews identify heavy reliance on BBOB, heterogeneous feature protocols, and limited cross-benchmark evidence as continuing concerns [2]. Random instance partitions and within-training scores are therefore insufficient evidence for transfer to unseen function families or benchmark suites.

This paper studies a decision that precedes feature acquisition: given an optimization prefix and its algorithm-agnostic behavioral history, should a fixed landscape-descriptor query be executed at the current state? We call this setting *Decision-before-Feature*. It differs from trajectory-based per-run selection, which constructs a trajectory representation and chooses a second-stage algorithm [11, 14], and from dynamic algorithm selection,

ORCID(s):

which repeatedly changes search actions during a run [6]. It also differs from landscape-aware operator or population control [26, 28]. The object here is not a parameter, operator, or optimizer update rule, but whether to acquire additional information before a downstream portfolio decision.

We formulate query value as the normalized performance difference between a native no-query continuation and a query-plus-selector path from the same complete optimizer state, less a weighted time cost. Memory is measured and reported separately in the primary analysis, for which $\lambda_M = 0$. The query evaluations already reduce the remaining optimization budget and are not deducted twice. Offline labels use completed candidate continuations, whereas the deployed gate observes only behavior available before the query. Query descriptors, function identity, algorithm identity, and optimizer-specific parameters are excluded. The resulting question is not whether landscape features can sometimes support selection, but when the expected downstream difference of the evaluated query is sufficient to justify its acquisition cost.

The study is organized around five research questions. RQ1 characterizes the statewise utility of the primary descriptor_cheap query. RQ2 tests whether algorithm-agnostic behavior adds decision value beyond elapsed budget alone. RQ3 compares the resulting policy with fixed, random, traditional landscape-based, and portfolio references under common resource accounting. RQ4 tests the frozen procedure on held-out BBOB families, CEC2017, CEC2022, and engineering problems, with each suite reported separately. RQ5 examines the prespecified T0–B3 progression, model-appropriate coefficient or discriminant stability, maturity–utility associations, and variation across family, dimension, and query configuration. RQ4 and RQ5 remain empirical questions; their intended evaluations are not evidence that transfer or explanatory stability has already been obtained.

The contributions are methodological. First, we define execution of a fixed landscape-descriptor query as a resource-aware prediction problem with a state-conditional net-utility target. Second, we construct the gate input from permutation-invariant native-update histories while excluding algorithm identity and query descriptors. Third, we separate this gate from a query-specific downstream action-loss selector and make native continuation and population-transfer initialization explicit. Fourth, we specify nested function-family out-of-fold model selection and train-only threshold fitting, together with an evaluation structure connecting decision utility, final optimization loss, query use, and measured resource cost. Among the literature reviewed here, we did not identify the same combination of pre-query information timing, matched-state utility labels, and an independently acquired query. Whether the method yields a predictive or resource advantage is determined only by RQ1–RQ5.

The remainder of the paper reviews landscape-based and trajectory-based selection, behavior characterization, adaptive optimization, and generalization; formalizes the decision and utility target; describes the offline construction and learning procedure; and presents the RQ-aligned evaluation and reproducibility design.

2. Related work

2.1. Landscape representations and automated algorithm selection

Fitness-landscape analysis describes properties of an optimization problem that may be relevant to search, including ruggedness, modality, neutrality, fitness distributions, and variable interactions. The corresponding measures differ in their assumptions, sampling procedures, dependence on a search process, and computational demands [16]. ELA operationalizes this idea through numerical descriptors computed from sampled objective values. Its original formulation proposed relatively inexpensive, computer-generated features as a route from problem characterization to automatic algorithm selection [18]. Software such as pflacco now exposes multiple feature families with different sample requirements and domains of applicability [22]. Estimated ELA values can also depend materially on the sampling strategy, and a mismatch between the designs used for training and deployment can degrade the resulting predictor [23]. The sample design is therefore part of the representation, not a neutral implementation detail.

In automated algorithm selection (AAS), a conventional pipeline computes an instance representation and predicts which portfolio member should be used. ELA combined with supervised learning has been evaluated against the single best solver (SBS) and virtual best solver (VBS), often under leave-one-function-out or related problem-oriented splits [12]. More recent systems retain the sample–feature–selector structure while changing the portfolio, feature set, or learning model [7]. These studies establish the downstream use of landscape information. They do not determine whether an independently sampled descriptor query should be acquired at a particular state before its features exist.

2.2. Feature acquisition as a resource cost

Landscape information is acquired rather than supplied. Objective evaluations used for feature sampling reduce the budget left for optimization, while feature computation adds computational demand. Reusing the sample as an optimizer population can reduce the effective FE charge in a particular pipeline [12], but this does not make the representation cost-free in general. Runtime and peak memory are therefore measured separately in the present study. A recent preprint study of feature budgets found that the preferred sampling fraction varied with the portfolio, problem set, dimension, target budget, and selection scenario [1]. Similarly, ELA-based selection experiments may restrict feature groups when the acquisition budget is limited [7].

These findings motivate cost-aware evaluation, but their decision variable is usually the portfolio action after features

have been obtained. In contrast, Decision-before-Feature compares executing and skipping the same prespecified query from a shared search state. Query FEs enter through the reduced continuation budget rather than a second penalty; normalized time is penalized explicitly, while memory is reported and examined only in sensitivity analysis. The target is therefore not a classifier for whether a feature family is “cheap.”

2.3. Trajectory-based and per-run algorithm selection

The closest antecedents replace an independent landscape sample with information collected during an initial optimizer run. Jankovic et al. begin with CMA-ES, compute trajectory-based landscape features, predict the performance of candidate second-stage optimizers, and warm-start the selected one [11]. In that formulation, the trajectory representation is constructed for every run at a fixed switch point; the decision is which second-stage optimizer to use, not whether to construct the representation. Kostovska et al. extend this per-run setting to compare trajectory ELA, time series of CMA-ES internal state variables, and their combination [14]. The latter study also reports that its BBOB findings did not transfer directly to YABBOB, showing that such transfer cannot be assumed; trajectory-coverage mismatch is discussed there as a possible explanation. Both internal evaluations retain the same BBOB functions across training and test partitions: the CEC study groups five-dimensional observations by instance identifier, while the PPSN study uses run- and instance-level partitions in five and ten dimensions. These protocols answer different questions from the present function-family OOF design; they are antecedents for representation and warm-starting, not precedents for the present evaluation design.

Probing trajectories provide another algorithm-centric representation. A short run records an objective-value sequence, and the resulting time series supports a later portfolio choice [24]. A subsequent benchmark compared multiple time-series classifiers under both leave-one-instance-out and the harder leave-one-problem-out evaluation [25]. This result is relevant here because it shows that the predictive model and problem partition remain part of a trajectory selector’s empirical specification. Nevertheless, probing is performed before every selection, the target is the best portfolio algorithm, and the output is an optimizer choice. The probe itself still consumes evaluations even though it is not a separate landscape sample. It does not construct paired skip/query utility for a separate landscape query or decide whether that query should be acquired.

2.4. Trajectory representations and behavior characterization

Other work studies what optimizer histories represent rather than which action to take. Search trajectory networks summarize visited regions and transitions for comparative analysis and visualization [19]. DynamoRep concatenates generation-wise coordinate and fitness summaries,

thereby preserving longitudinal population dynamics without a separate objective sample; Opt2Vec instead learns a representation from populations observed during a run [3, 13]. DynamoRep fits separate classifiers for each optimizer and evaluates three-dimensional BBOB class identification, while Opt2Vec likewise reports problem classification. These tasks neither establish a fixed-dimensional, algorithm-agnostic state representation nor predict the utility of acquiring a separate landscape query.

Behavioral analysis provides a complementary vocabulary. Hayward and Engelbrecht characterize exploration, exploitation, locality, communication, and evaluation effort with a suite of whole-run measures, explicitly treating behavior as problem dependent [8]. Some of those measures rely on a known optimum, persistent individual identity, or algorithm interaction structure and hence are unsuitable as pre-query inputs here. Mbasso et al. use distance-to-reference decay, directional entropy, and stagnation to diagnose late-stage behavior and trigger optimizer-specific interventions [17]. Only the distance-decay and stagnation semantics motivate the present state representation: distance is redefined relative to the current population best, while particle-wise directional entropy is excluded because it requires persistent cross-generation identities. The reported interventions do not validate the study-specific Search Maturity construction or establish that behavior predicts landscape-query utility. More generally, computability alone does not make a behavior indicator informative or complementary: sensitivity, redundancy, single-run computability, and the information available at deployment all constrain its use [9].

2.5. Adaptive search and dynamic algorithm selection

Landscape-aware adaptive methods change the search mechanism after observing landscape or performance information. Examples include mutation-operator selection in differential evolution and adaptation of subpopulation number and partitioning in artificial bee colony optimization [26, 28]. Adaptive landscape analysis instead studies how locally sampled landscape descriptors evolve during an optimizer run, with the longer-term aim of online selection or configuration [10]. It therefore computes descriptors on 2,000 additional points sampled from the current CMA-ES distribution at prescribed target levels. “Cheap” denotes that the chosen descriptor families need no further evaluations beyond that sample, not that acquisition is free. The preliminary study neither makes a pre-descriptor acquisition decision nor couples the descriptors to a paired skip/query utility. Dynamic algorithm selection can make repeated choices during a run; RL-DAS, for example, represents the process as a reinforcement-learning problem over differential-evolution variants using landscape and algorithm-history information [6]. Dynamic algorithm configuration makes the control problem explicit through a state space, an action space over hyperparameter settings or components, a reward,

Table 1

Closest task formulations compared by decision object, information time, acquisition cost, and state transition. “Extra FEs” denotes an independent landscape sample beyond an existing prefix or probe.

Approach	Decision target	Information at decision	Extra FEs	Frequency and transition
Static ELA-based AAS [7, 12]	Portfolio algorithm	Acquired landscape descriptors	Usually yes	Once, before the chosen optimizer starts
Jankovic et al. [11]	Second-stage algorithm	ELA computed from a fixed CMA-ES prefix	No separate sample	Once, with algorithm-specific warm-starting
Kostovska et al. [14]	Second-stage algorithm	Trajectory ELA and/or CMA-ES state series	No separate sample	Once, with algorithm-specific warm-starting
Probing trajectories [24, 25]	Portfolio algorithm	Objective-value sequence from a short probe	No separate landscape sample	Once, after probing
RL-DAS [6]	Differential-evolution variant	Landscape, search history, and context	Not isolated as a fixed query	Repeated online scheduling
DACBench [5]	Hyperparameter or algorithm component	Target-algorithm state and instance context	Not a landscape-query decision	Repeated dynamic configuration
Decision-before-Feature	Execute or skip one fixed query	Pre-query algorithm-agnostic behavior	Only when called	Opportunities until one call; native continuation or population transfer

and a decision frequency; DACBench includes both CMA-ES step-size control and repeated selection among modular evolutionary-algorithm components [5]. Parameter control and MetaBBO cover still broader choices over parameters, components, structures, and learned update rules [15, 27]. The broader adaptive-operator-selection literature likewise separates stateless credit-based rules from state-based policies that map solution, problem, or optimization-process features to an operator [20]. In both cases, the selected action is a search operator; the acquisition of an independent descriptor query is not the action being evaluated.

Those approaches optimize a search action or solver mechanism. The present method uses offline supervised learning and leaves the portfolio algorithms unchanged. Its gate decides whether to acquire information; only after a call does a separate, fixed action-loss selector choose between native continuation and population-transfer initialization. Thus, its action is neither an optimizer hyperparameter nor a repeated component choice. Table 1 summarizes this distinction for the closest task formulations.

2.6. Evaluation splits and the remaining gap

Representation quality and selector performance depend strongly on how problems are divided between training and evaluation. In an affine-BBOB study, several landscape representations that performed well under easier splits did not outperform SBS under the most demanding problem-oriented splits [4]. A recent representation survey likewise highlights inconsistent sampling and split protocols, heavy dependence on BBOB, and limited cross-benchmark evidence [2]. These observations motivate function-family out-of-fold selection and a held-out evaluation; they do not guarantee that the resulting gate transfers. Benchmarking work

further shows that leave-one-instance-out evaluation can reward representations that recover a problem class shared by related training and test instances without learning to rank algorithms on unseen problems, and that scale-sensitive prediction errors need not measure selection quality [21]. The present study therefore uses function-family partitions and selects the Decision Model by policy utility rather than raw action-loss regression error. This design addresses those two failure modes without establishing generalization in advance.

Among the literature reviewed here, we did not identify a method that makes the same decision at the same information time: before descriptors from a prespecified independent query exist, use only algorithm-agnostic behavior to predict whether the query’s downstream performance difference justifies its FE and computational cost. Decision-before-Feature addresses this corpus-scoped gap with matched-state offline labels, an explicit separation between the query gate and the downstream selector, a time-only comparator, and nested function-family evaluation. The literature supports this formulation and its design constraints; it does not supply evidence for the pending RQ1–RQ5 outcomes or for held-out and cross-benchmark generalization.

3. Problem formulation

3.1. State-conditional analysis selection

Let $p \in \mathcal{P}$ denote a bounded continuous black-box minimization problem with objective

$$\min_{x \in \mathcal{X}_p} f_p(x), \quad \mathcal{X}_p = [\ell_p, u_p] \subset \mathbb{R}^d. \quad (1)$$

The objective is available only through function evaluations (FEs). A run has a fixed total budget $B = FE_{\text{total}}$ and uses a

portfolio

$$\mathcal{A} = \{\text{DE, PSO, CMA-ES, SHADE}\}. \quad (2)$$

At a decision opportunity t , the native optimizer state is

$$s_t = (X_t, y_t, H_t, \omega_t, e_t), \quad (3)$$

where X_t and y_t are the current population and its objective values, H_t is the native-update history required by the behavior representation, ω_t collects the optimizer's internal dynamic variables and random-number-generator state, and e_t is the number of FEs already consumed. The complete state in Eq. (3), rather than a reconstructed population-only approximation, is the common starting point for all continuations.

The analysis-selection decision is binary, $d_t \in \{0, 1\}$, with

$$d_t = \begin{cases} 0, & \text{skip query } q, \\ 1, & \text{execute } q \text{ and invoke its selector.} \end{cases} \quad (4)$$

The primary analysis evaluates one query, `descriptor_cheap`: a fixed 16-dimensional low-cost landscape descriptor computed from a Latin-hypercube sample of size $50d$. It consumes $FE_q = 50d = 0.05B$ objective evaluations. Two additional, prespecified query configurations examine representation and cost dependence. Consequently, d_t does not mean ‘‘perform arbitrary landscape analysis’’, and conclusions for the primary query do not automatically extend to other ELA methods, all `pflacco` feature groups, or learned representations.

3.2. No-query and query paths

The default optimizer $a_0 \in \mathcal{A}$ is the single best solver (SBS) determined exclusively from complete-budget BBOB training outcomes. The main protocol uses the same algorithm for the prefix and the default path. If the query is skipped, the complete checkpoint state is continued natively for

$$B_t^{\text{skip}} = B - e_t \quad (5)$$

additional FEs. Let $P_{\text{skip}}(s_t)$ denote the resulting final loss, with smaller values preferred.

If the query is executed, its sample is evaluated and its fixed descriptors are passed to a downstream, query-specific selection reference. The unique action set at a state whose prefix algorithm is a_t is

$$\mathcal{A}(s_t) = \{\text{continue_current}\} \cup (\mathcal{A} \setminus \{a_t\}). \quad (6)$$

Thus, the four-algorithm portfolio always yields four non-duplicated actions. The `continue_current` action preserves ω_t and continues the native state. Each of the other three actions initializes the target optimizer once from the checkpoint population, objective values, and best-so-far position; the target optimizer receives newly initialized internal variables. This transition is termed *population-transfer initialization*.

Query sample points are not inserted into the continuation population. This explicit distinction is motivated by trajectory-based warm-starting studies, where the representation, switch point, and state-transfer rule jointly define the second-stage selection problem [11, 14]. Population-transfer initialization here is a fixed study design and is not assumed equivalent to either paper's optimizer-specific warm start.

Every query action receives the same remaining optimization budget

$$B_t^q = B - e_t - FE_q. \quad (7)$$

Let $L(s_t, a)$ be the observed final loss after taking action $a \in \mathcal{A}(s_t)$ under Eq. (7). The fitted selection reference chooses $\hat{a}_t$, and the observed query-path loss is

$$p_q(s_t) = L(s_t, \hat{a}_t). \quad (8)$$

Equations (5) and (7) enforce an equal total FE budget. The FEs used by the query already reduce the budget available to the query continuation and must not be charged a second time in the utility.

3.3. Observed performance difference and net utility

For a minimization problem, the raw observed performance difference is

$$G_{\text{raw}}(s_t) = P_{\text{skip}}(s_t) - p_q(s_t). \quad (9)$$

Accordingly, $G_{\text{raw}} > 0$ means that the query path produced the smaller final loss, whereas $G_{\text{raw}} < 0$ means that the no-query path produced the smaller final loss. To reduce the influence of objective offsets and scales across problems, the main label uses

$$\kappa(s_t) = \max\left(\left|P_{\text{skip}}(s_t)\right|, \left|p_q(s_t)\right|, 10^{-12}\right), \quad (10)$$

$$G(s_t) = \frac{G_{\text{raw}}(s_t)}{\kappa(s_t)}.$$

The normalized time cost includes query sampling and feature-computation time, together with online selection time, relative to the no-query continuation time:

$$C_T(s_t) = \frac{T_q(s_t) + T_{\text{sel}}(s_t)}{\max(T_{\text{skip,opt}}(s_t), 10^{-12})}. \quad (11)$$

The primary study reports memory separately rather than including it in the label and therefore fixes $C_M(s_t) = 0$. The net utility of executing query q is

$$U_q(s_t) = G(s_t) - \lambda_T C_T(s_t) - \lambda_M C_M(s_t). \quad (12)$$

The primary target uses $\lambda_T = 1$ and $\lambda_M = 0$. Values $\lambda_T \in \{0, 0.25, 0.5, 1, 2\}$ are retained for a prespecified sensitivity analysis, but they do not replace the primary target. Population-transfer effects are already present in $L(s_t, a)$ and likewise are not subtracted again.

The continuous value in Eq. (12) is retained. Its associated binary training label is

$$y_q(s_t) = \mathbf{1}[U_q(s_t) > 0]. \quad (13)$$

For deployment, a decision model produces a continuous score $z(s_t)$ from information available before the query, and the policy calls the query when

$$\hat{d}_t = \mathbf{1}[z(s_t) > \tau_{\text{OOF}}]. \quad (14)$$

The threshold τ_{OOF} is estimated only from BBOB-training function-family out-of-fold predictions; it is not tuned on BBOB validation or an external benchmark. For Ridge, z is a predicted continuous utility. For LDA and logistic regression, it is a classification score for Eq. (13); such a score is not interpreted as a calibrated utility magnitude.

3.4. Diagnostic decomposition and claim boundary

For label diagnosis only, define the best observed action among the four actions actually run from the shared state,

$$L_{\text{best}}(s_t) = \min_{a \in \mathcal{A}(s_t)} L(s_t, a). \quad (15)$$

The raw performance difference then has the identity

$$V_{\text{potential}}(s_t) = P_{\text{skip}}(s_t) - L_{\text{best}}(s_t), \quad (16)$$

$$R_{\text{selector}}(s_t) = p_q(s_t) - L_{\text{best}}(s_t), \quad (17)$$

$$G_{\text{raw}}(s_t) = V_{\text{potential}}(s_t) - R_{\text{selector}}(s_t). \quad (18)$$

The minimum in Eq. (15) is called the *best observed action*. It is neither the static virtual best solver nor a deployable decision rule, because it uses losses revealed only after all candidate continuations have been run. The decomposition is descriptive: shared-state action runs do not by themselves identify a causal effect of handoff or of any individual behavior variable.

The resulting research estimand is state- and query-specific: the net value of invoking a fixed query and its fitted downstream selector instead of natively continuing the training-derived SBS from the same complete state. Equations (4)–(18) specify the estimand and label construction. They do not imply that the primary query is avoidable in most states, that behavior predicts its utility, or that the learned policy transfers to an external benchmark; each proposition requires its corresponding empirical evaluation.

4. Decision-before-Feature method

4.1. Overview

Decision-before-Feature is trained offline. Native optimizer trajectories are first collected under a fixed sampling protocol. Each accepted state is converted into a permutation-invariant behavior vector available before landscape descriptors are computed. From the same complete state, the no-query continuation and all four query-adjusted actions are then run to generate observed losses. A fixed, query-specific action-loss regressor produces the

downstream selected action, after which Eq. (12) yields a supervised utility label. Finally, a lightweight decision model maps behavior alone to a query score. The deployed controller performs no online parameter update.

Figure 1 summarizes the separation between offline data construction, query-specific downstream selection, decision-model freezing, and online execution. The use of trajectory-derived summaries follows the general observation that optimizer histories carry information about search dynamics and problem–algorithm interaction [2, 13, 19]; their value for the present query decision remains an empirical question.

This separation gives the two supervised components different roles. The *selection reference* operates after query acquisition and may use the query descriptors. The *decision model* operates before query acquisition and is prohibited from using query descriptors, function identity, dimension, algorithm identity, optimizer-specific parameters, action losses, or best-observed-action diagnostics.

4.2. Native-state trajectory and window alignment

For every portfolio optimizer, a complete state contains its population, objective values, best-so-far solution, generation or native-update count, internal dynamic variables, and random-number-generator state. Restoring this state continues the same optimizer rather than constructing a new run from population values. This distinction is essential for PSO velocities, CMA-ES distribution parameters, and the adaptation histories of DE and SHADE.

Let e_t be the actual FE count at an accepted state and let $w \in \{0.02, 0.05, 0.10\}$ be a nominal behavior window. Among the complete native updates preceding t , the anchor is

$$a(t, w) = \arg \max_u \{e_u : e_u \leq e_t - \text{round}(wB)\}. \quad (19)$$

For population size $N = 40$, the realized span satisfies

$$\text{round}(wB) \leq e_t - e_{a(t,w)} < \text{round}(wB) + N. \quad (20)$$

Rates and slopes use the realized FE-ratio span, not the nominal value. The realized ratio, FE count, and number of native updates are retained as measurement metadata and are excluded from the decision input.

Behavior vectors are evaluated only at accepted complete-update states. The budget milestones, behavior-triggered opportunities, rearming rules, and per-phase quotas are specified in Section 5; opportunity labels determine state sampling but are excluded from the Decision Model.

4.3. Permutation-invariant behavior representation

For distance-based population summaries, coordinates are mapped to the unit hypercube using the problem bounds,

$$\tilde{x}_{ij} = \frac{x_{ij} - \ell_{pj}}{u_{pj} - \ell_{pj}}. \quad (21)$$




Figure 1: Decision-before-Feature protocol. The upper panel constructs state-aligned utility labels and freezes the downstream Selection Reference, Decision Model, preprocessing, and threshold using BBOB training families. The lower panel applies those frozen components at an online decision opportunity. Query descriptors enter only the downstream Selection Reference; the pre-query Decision Model receives algorithm-agnostic behavior.

This unit-cube convention defines distance normalization. No correspondence between row i at two checkpoints is assumed. The current population diversity is the average pairwise distance

$$D_t = \frac{2}{N(N-1)} \sum_{1 \leq i < j \leq N} \|\tilde{x}_i^{(t)} - \tilde{x}_j^{(t)}\|_2. \quad (22)$$

For an anchor population and the current population, the empirical assignment distance is

$$W_t^{(w)} = \frac{1}{N} \min_{\pi \in S_N} \sum_{i=1}^N \|\tilde{x}_i^{(a)} - \tilde{x}_{\pi(i)}^{(t)}\|_2, \quad (23)$$

and its rate divides Eq. (23) by $(e_t - e_a)/B$. Centroid shift and its rate are computed analogously. Centroid-shift coherence is the ratio of centroid displacement to $W_t^{(w)}$, clipped to $[0, 1]$ and set to zero when the population does not change. Covariance trace and effective-rank changes use the unit-cube populations. Current covariance spectral concentration is computed from the raw-coordinate covariance. For its nonzero eigenvalues λ_j and $r = \min(d, N-1)$,

$$C_t = \frac{r \sum_j \lambda_j^2 / (\sum_j \lambda_j)^2 - 1}{r - 1}, \quad (24)$$

with $C_t = 0$ for a degenerate covariance or $r \leq 1$.

Fitness changes use a shift-invariant scale. Let

$$I_0 = \max(\text{IQR}(y_{\text{init}}), 10^{-12}), \quad (25)$$

where y_{init} is the evaluated population immediately after optimizer initialization and before any native update. Improvement rate over a short window is

$$IR_t^{(w)} = \frac{f_{\text{best}}(a) - f_{\text{best}}(t)}{I_0 (e_t - e_a) / B}. \quad (26)$$

Improvement frequency is the proportion of native-update intervals in which best-so-far fitness decreases strictly. To compare fitness distributions without individual identities, the two endpoint vectors are sorted. If $y_{(i)}^{(a)}$ and $y_{(i)}^{(t)}$ denote their empirical quantiles, the method records the fraction for which $y_{(i)}^{(a)} - y_{(i)}^{(t)}$ exceeds the numerical strict-improvement tolerance, their mean signed improvement divided by I_0 and the realized ratio span, and their mean absolute difference on the same scale.

Stagnation is based on the last strict best-so-far improvement. With e_{last} denoting its FE count,

$$S_t^{(0.10)} = \frac{\min\{\max[(e_t - e_{\text{last}})/B, 0], 0.10\}}{0.10}. \quad (27)$$

Distance decay compares the population's mean unit-cube distance to its current population-best at the long-window endpoints. Convergence rate is the relative decrease in D_t divided by the realized FE-ratio span. The reference point in these calculations is the current population-best, not a known global optimum.

The representation borrows the general idea that longitudinal population summaries can characterize problem–algorithm interaction [3], but uses one fixed-dimensional definition across all four optimizers and predicts query utility rather than a problem class. Distance-decay and stagnation

Table 2

Frozen behavior groups used in the feature-group comparison.

Group	Count	Added information	Role
T0	1	Actual FE ratio only	Time-only controller
B1	19	Progress, diversity, set and fitness-distribution change, stagnation, convergence, and elite concentration	Core permutation-invariant behavior
B2	25	B1 plus fitness-spread slope, population and elite centroid shifts, covariance trace ratio, covariance effective rank, and diversity recovery	Longitudinal set dynamics
B3	31	B2 plus set-motion proxies and three maturity variables	Primary full input

semantics are also informed by behavior-characterization studies [8, 17]. Measures requiring a known optimum, persistent particle identities, or an optimizer-specific interaction graph are excluded; in particular, particle-wise directional entropy is not an input.

The behavior representation has 34 unique outputs: 31 model-eligible variables and three diagnostic-only variables. The nested groups are summarized in Table 2.

The raw fitness standard deviation, population-overlap statistic, and best-distance–fitness correlation are diagnostic only. In particular, the last quantity is close to a local landscape proxy and is excluded from the pre-query model.

4.4. Operational search-maturity variables

Search Maturity is an operational, study-specific summary rather than an established convergence measure. Define

$$h_+(z) = \frac{\max(z, 0)}{1 + \max(z, 0)}, \quad h_-(z) = \frac{1}{1 + \max(z, 0)}, \quad (28)$$

$$c(z) = \min\{\max(z, 0), 1\}.$$

Using the defined behavior variables, exploration stabilization is the mean of

$$h_+(-\dot{D}_t), \quad h_-(\dot{W}_t), \quad c(C_t), \quad c(K_t), \quad (29)$$

where $\dot{D}_t$ is the medium-window diversity slope, $\dot{W}_t$ is the population assignment-distance rate, and K_t is centroid-shift coherence. Exploitation saturation is the mean of

$$c(S_t), \quad h_-(IR_t), \quad h_+(DD_t), \quad h_+(CR_t), \quad h_-(E_t), \quad (30)$$

where DD_t , CR_t , and E_t denote distance decay, convergence rate, and elite-to-population diversity ratio. The primary maturity variable is

$$M_t = ES_t(1 - XS_t). \quad (31)$$

The B3 group also contains the bounded linear alternative $c[(ES_t + 1 - XS_t)/2]$ and an exploration-to-exploitation ratio. Undefined components remain missing and are handled by training-only median imputation. Equation (31) is a hypothesis-bearing representation whose utility must be established by the T0–B3 comparison and external evaluation; it is not treated as an independently validated causal mechanism.

4.5. Query-specific downstream selection reference

For every eligible shared state, all actions in Eq. (6) are run under the query-adjusted budget. Because raw objective values vary across problems, each action target is normalized within its state:

$$\tilde{L}(s_t, a) = \frac{L(s_t, a) - \min_b L(s_t, b)}{\max(\max_b L(s_t, b) - \min_b L(s_t, b), 10^{-12})}. \quad (32)$$

The selection reference predicts the four targets jointly from the 31 B3 behavior variables, the fixed descriptors of the current query, and the continuous remaining-budget ratio $(B - e_t - FE_q)/B$:

$$\hat{L}(s_t) = g\left(b_t, \phi_q(p), \frac{B - e_t - FE_q}{B}\right), \quad (33)$$

$$\hat{a}_t = \arg \min_a \hat{L}(s_t, a). \quad (34)$$

The fixed estimator is a multi-output random-forest regressor with 200 trees, minimum leaf size two, median imputation, and random seed 1701. BBOB-training rows receive five-fold function-family cross-fitted predictions so that a training utility label never uses a selector fitted on the same family. The final selector is refitted on all BBOB-training families for BBOB validation and external evaluation. This selection reference is a downstream experimental component, not the proposed contribution.

For each selected action, three relations distinguish whether it equals the default, whether it equals the prefix, and whether a handoff is required. Writing h_t for the handoff indicator, i_t for equality with the prefix, and r_t for the transition type, these relations satisfy

$$h_t = \neg i_t^{\text{prefix}} = \mathbf{1}[r_t = \text{population transfer initialization}], \quad (35)$$

Here h_t , i_t^{prefix} , and r_t correspond, respectively, to the reported handoff indicator, equality with the prefix action, and transition type. Equality with the default action is reported independently. These relations hold by definition for every state, support stratified reporting, and are not Decision Model inputs.

4.6. Decision-model selection and threshold fitting

The active candidates are fixed to (i) linear discriminant analysis with default estimator settings, (ii) logistic regression with $C = 1$, balanced class weights, and the L-BFGS solver, and (iii) Ridge regression with $\alpha = 1$. Each candidate is a pipeline containing training-fold median imputation and standard scaling. LDA and logistic regression fit Eq. (13); Ridge fits the continuous primary utility. These three candidates are compared only with the frozen B3 input group.

Model selection is performed once and is nested entirely within BBOB training. The outer loop is a five-fold GroupK-Fold over function families. Within each outer fit partition, a four-fold family out-of-fold procedure supplies scores used to fit an outer-fold decision threshold. Candidate quality is the mean decision utility obtained after concatenating all outer holdouts. Ties are resolved by the prespecified order LDA, logistic regression, then Ridge. Neither T0, B1, B2 nor BBOB validation contributes to this model-family choice.

After the B3 comparison freezes the model family, that same named estimator family is used for T0, B1, B2, and B3 with identical samples, family folds, training-fold preprocessing, and threshold-fitting procedure. Each input group fits its threshold only from its own BBOB-training family-OOF scores. The RQ5 feature-group analysis reports all four prespecified groups; it neither repeats model-family selection nor selects a feature group from their observed results. B3 remains the primary deployable representation.

After a model family is selected, a separate five-fold function-family out-of-fold pass over all BBOB-training rows produces one score per row from a model that did not see that row's family. The deployable threshold is

$$\tau_{\text{OOF}} \in \arg \max_{\tau} \sum_i \mathbf{1}[z_i^{\text{OOF}} > \tau] U_{q,i}. \quad (36)$$

If thresholds yield the same utility sum, the threshold calling fewer rows is preferred. The selected estimator is then fitted on all BBOB-training families and deployed with the fixed threshold. BBOB validation is evaluation-only. The optional score-neighborhood review uses the tenth percentile of $|z_i^{\text{OOF}} - \tau_{\text{OOF}}|$ from BBOB training; it is performed only after model and threshold freezing and gives all comparison policies the same additional opportunities.

At deployment, the default SBS advances natively while the streaming behavior extractor observes complete updates. At each accepted decision opportunity, the controller evaluates Eq. (14). If the score does not cross the threshold, native optimization continues. If it crosses, the fixed query is executed once, the downstream selector chooses an action, and the selected continuation consumes the remaining query-adjusted budget. No subsequent online training or repeated query-type selection is performed.

5. Experimental setup

5.1. Research questions and evaluation design

The empirical study addresses the five prespecified questions in Table 3. Each question is linked to an estimand,

a comparison set, and uncertainty summaries that respect dependence among states from the same optimization trajectory and function family.

5.2. Internal benchmark and function-family partition

The internal benchmark is the noiseless BBOB suite. The training set comprises functions 1, 2, 3, 4, 6, 7, 8, 10, 11, 12, 15, 16, 17, 18, 20, 21, 22, and 23; functions 5, 9, 13, 14, 19, and 24 form the held-out validation set. Functions, rather than randomly selected transformed instances, define the grouping variable. This partition addresses evidence that instance-level splits can overstate algorithm-selection generalization [4]. Both sets use dimensions $d \in \{10, 20, 40\}$, instances 1, 2, 3, optimizer seeds $1, \dots, 30$, population size 40, and total budget $B = 1000d$. The portfolio contains DE, PSO, CMA-ES, and SHADE, giving 54 training and 18 validation function–dimension combinations.

The single best solver (SBS) is estimated exclusively from terminal outcomes at $FE = B$. For each problem–algorithm pair, final best fitness is first averaged over seeds. The four algorithms are then ranked within each problem, with smaller mean fitness preferred and average ranks assigned to ties. These ranks are averaged over the training problems; any remaining tie is resolved in the prespecified order DE, PSO, CMA-ES, SHADE. Intermediate decision states, including the last state at or before approximately $0.60B$, do not contribute to the SBS estimate.

BBOB validation is a held-out internal evaluation set. It remains untouched during query selection, preprocessing, model-family selection, input-group comparison, threshold fitting, and score-neighborhood definition.

5.3. External benchmarks

External evaluation is organized as non-pooled, suite-specific analyses. The CEC2017 analysis comprises functions 1–29, dimensions 10, 30, and 50, seeds $1, \dots, 30$, population size 40, and budget $B = 1000d$. The four-algorithm portfolio and dynamic decision-opportunity protocol are unchanged, and the controller, downstream selector, imputation values, and threshold are held fixed at their BBOB-training estimates.

CEC2022 and engineering problems form separate components of RQ4. Their problem sets, dimensions, bounds, budgets, success targets, and failure rules are specified independently of policy-outcome inspection. For every external suite, query-extraction failures remain in the coverage denominator rather than being removed from the analysis. Missing descriptors may be imputed only with BBOB-training medians. Held-out BBOB, CEC2017, CEC2022, and engineering estimates are reported separately; neither partial coverage nor agreement on a subset of suites is interpreted as unconditional cross-benchmark transfer.

Table 3

Research questions and corresponding evaluation criteria.

RQ	Question	Evaluation criteria
RQ1	Statewise value of the primary query	Utility distribution and decomposition; proportion with $U_q \leq 0$; family-balanced effects and intervals.
RQ2	Information beyond elapsed budget	Nested training-family OOF selection; paired B3–T0 effects; frozen BBOB-validation evaluation.
RQ3	End-to-end policy trade-off	Call quality, final performance, FE use, runtime, and paired comparisons with all required policies.
RQ4	Transport of the frozen procedure	Separate held-out BBOB, CEC2017, CEC2022, and engineering estimates, with coverage and query failures.
RQ5	Behavioral associations with query value	T0–B3 increments, model-parameter stability, maturity–utility association, and prespecified strata.

Table 4

Prespecified landscape-descriptor queries.

Query	Sample	FE share	Features	Role
descriptor_cheap	LHS 50d	5%	16	Primary
pflacco_standard	Same LHS 50d	5%	37	Configuration robustness
pflacco_broad	Independent LHS 100d	10%	52	Configuration robustness

5.4. Landscape-query configurations

Table 4 defines the three prespecified queries. Primary inference concerns `descriptor_cheap`; the two `pflacco` configurations evaluate dependence on representation and acquisition cost.

The cheap and standard queries use the same evaluated (X, y) sample, yielding a paired comparison of descriptor representations at a common FE cost. The broad query uses an independent 100d design. The standard configuration comprises PCA, nearest-better clustering, dispersion, information-content, and ELA-distribution groups. The broad configuration adds ELA level-set and sample-derived fitness–distance-correlation groups, while excluding the full quadratic meta-model group, cell mapping, and feature groups that require additional objective evaluations. Query sample points are not inserted into the optimizer population. Each query defines a separate downstream selector, utility target, decision model, threshold, baseline comparison, and external evaluation; observations from different query identities or FE designs are not pooled.

5.5. Dynamic decision opportunities

Offline label construction and policy evaluation use the same dynamic budget–event sampling scheme. Complete native updates are monitored when they cross a grid from $0.20B$ to $0.60B$ in increments of $0.01B$. The mandatory budget milestones are 0.20, 0.22, 0.24, 0.26, 0.28, 0.30, 0.34, 0.38, 0.42, 0.46, 0.50, and 0.60. Five event types may add decision states: resumed strict improvement; a stagnation span reaching $0.05B$; an absolute medium-window covariance-effective-rank change reaching 0.20; unit-cube elite-centroid displacement divided by $\sqrt{d}$ reaching 0.05; and relative diversity recovery reaching 0.10. Covariance-effective-rank-change, elite-migration, and diversity-recovery events rearm

when their respective quantities fall below 0.10, 0.025, and 0.05. Stagnation rearms after a new strict best-so-far improvement, whereas resumed improvement rearms when short-window improvement frequency returns to zero.

Each complete update that crosses one or more monitor points is evaluated once. If a crossing includes a milestone, the event is associated with that milestone and does not consume an event-only quota or spacing anchor. Otherwise, the most recent crossed monitor point supplies the nominal location. Each of the early, middle, and late phases admits at most two event-only states, with at least $0.02B$ separation in realized FE ratio. Suppressed crossings still update event rearming. This scheme admits 12–18 states per run, each with unit sample weight. The realized FE/B represents elapsed budget and is the sole T0 covariate; exact integer FE aligns shared-state outcomes. Nominal milestones and native-update counts describe the sampling process but do not add Decision Model inputs.

5.6. Matched states and analysis population

The analysis pairs all alternatives at an identical complete native state. For each eligible state, it combines permutation-invariant behavior, the shared 50d and independent 100d LHS samples, the three descriptor representations, and the query-budget-specific outcome of every eligible action. A trajectory enters an estimand only when both its complete native history and terminal outcome satisfy the same state definition. Query-specific selection references, utilities, decision models, and policies are estimated from these matched observations under the family partitions described above.

The primary Decision Model population contains states for which both the prefix and default algorithm equal the

training SBS and the no-query path continues that optimizer. States generated from the other prefix algorithms are reserved for cross-prefix robustness, leave-one-prefix-algorithm-out analysis, and tests of portability across the four search processes. They are excluded from primary model fitting, threshold fitting, and primary result aggregation.

5.7. Comparison policies

Eight prespecified comparison roles are retained because they answer different benchmarking questions.

- **Never Query:** continue the complete SBS state natively and incur no query acquisition cost.
- **Always Query:** execute the fixed query at the first accepted decision opportunity and invoke the same downstream selector used by the proposed policy.
- **Random Analysis:** at each shared opportunity before a call, trigger with probability 0.5. Thirty independent repetitions use explicit integer random streams.
- **Traditional AAS:** execute the fixed query and apply the same statewise downstream selector, without a preceding Decision-before-Feature controller, consistent with the sample–feature–selector structure of landscape-based selection [7]. Under this definition, it induces the same first-opportunity decision as Always Query.
- **SBS:** use the training-derived static default defined above. In the primary population, the prefix and default are this same SBS, so its native complete-budget path coincides with Never Query; the separate name preserves its role as the standard algorithm-selection reference rather than introducing another online trajectory.
- **VBS:** use the static, per-problem hindsight choice among complete-budget portfolio outcomes. It is an upper reference rather than a deployable method and is distinct from the statewise best observed action.
- **Time-only Controller:** use the selected Decision Model family and the same training and threshold procedure as the full controller, with input restricted to $\{FE/B\}$.
- **Decision-before-Feature:** use B3 behavior with the selected model and the threshold estimated from full-training OOF predictions.

In the primary population, Never Query and SBS yield the same native terminal outcome, whereas Always Query and Traditional AAS yield the same first-opportunity query outcome. RQ3 therefore reports these roles as *Never Query / SBS* and *Always Query / Traditional AAS*; each shared outcome enters summaries and inferential comparisons once. All online policies share the same accepted decision opportunities, total FE budget, portfolio, query definition,

downstream selector, and population-transfer rule. State-only, query-only, and combined statewise selectors are additionally compared as diagnostics of the downstream selector; they do not replace the required policy baselines.

5.8. Evaluation measures

Decision-model family selection uses nested function-family OOF mean decision utility. AUROC, average precision, and Spearman correlation are secondary score-level measures. Continuous-utility RMSE is reported only for Ridge; classification probabilities and discriminant scores are not compared with continuous utility using RMSE.

Policy-level reporting includes query-call rate, mean decision utility, the proportion of called states for which $U_q > 0$, the fraction of positive-utility states called, and

$$\bar{U}_{\text{decision}}(\pi) = \frac{1}{n} \sum_{i=1}^n d_{\pi,i} U_q(s_i), \quad (37)$$

where $d_{\pi,i}$ is the call indicator of query-decision policy π . A skipped state contributes zero, so the Never Query role has zero decision utility by definition. The coincident SBS reference role and VBS are not query-decision policies and therefore remain not applicable for this quantity; mean observed utility conditional on calling is reported separately. Positive-utility capture is summarized by

$$\text{utility capture} = \frac{\sum_i U_{q,i} \mathbf{1}[U_{q,i} > 0] \mathbf{1}[\hat{d}_i = 1]}{\sum_i U_{q,i} \mathbf{1}[U_{q,i} > 0]}. \quad (38)$$

The positive-state capture rate and utility capture in Eq. (38) are reported separately. Calls with $U_q \leq 0$ are summarized by their count, their proportion among calls, and their accumulated utility loss.

End-to-end endpoints comprise final best objective value or error, target-based success rate and expected running time (ERT), total and query FEs, wall-clock time, behavior-extraction time, controller-inference time, query sampling and descriptor-computation time, downstream-selection latency, and peak memory. Success targets are suite specific and are fixed independently of the observed policy contrasts. Cost–performance reporting retains all eight prespecified comparison roles but uses six distinct primary-population outcomes: five online trajectories, including the two slash-labelled identities, and the VBS hindsight reference.

5.9. Statistical analysis

Repeated states from one trajectory are not treated as independent replicates. Descriptive summaries include state-level micro-averages and macro-averages balanced across function families. Results are also stratified by family, dimension, phase, and transition type. Support for the RQ1 statement that a majority of evaluated states do not benefit from the primary query requires the 95% lower confidence bound for the relevant proportion to exceed 0.5. Paired policy comparisons use nonparametric procedures, including Wilcoxon signed-rank comparisons and Friedman comparisons with Holm-adjusted follow-ups, at a prospectively

Table 5

State-level utility summary for the primary low-cost descriptor query. Utility is larger-is-better, whereas final action loss is smaller-is-better. Function families, rather than individual decision states, form the independent units for uncertainty quantification.

Analysis stratum	Coverage	States	$\Pr(U_q \leq 0)$	Mean U_q	Median U_q	Performance difference	Selector regret and query cost
Overall	All prespecified BBOB-train families, dimensions, and decision opportunities	–	–	–	–	–	–
By split	Separate BBOB-train family-OOF and frozen BBOB-validation estimates	–	–	–	–	–	–
By dimension	10D, 20D, and 40D	–	–	–	–	–	–
By decision stage	Each prespecified dynamic decision opportunity	–	–	–	–	–	–

TBD: numerical estimates and family-level intervals for all rows.

TBD
Distribution of U_q relative to zero, with function-family-aware uncertainty and separate views by dimension and actual FE ratio.

Figure 2: Distributional analysis of utility for the primary query. The zero reference distinguishes states in which the measured cost-adjusted utility is non-positive.

declared aggregation level. For every endpoint, analytically identical roles enter once; duplicated role labels add no ranks, paired contrasts, or multiplicity counts.

Equivalence is assessed only against a smallest effect of interest, margin, and primary procedure fixed before policy outcomes are examined. Confidence intervals use a pre-specified hierarchical resampling unit and replication count, and ERT uses suite-specific targets fixed independently of policy outcomes. Absence of statistical significance is not interpreted as equivalence or preservation of optimization performance.

6. Results

Results are structured around RQ1–RQ5. For each question, the corresponding estimand, family-level uncertainty, and policy or representation contrast are presented together; held-out BBOB and external-suite estimates are kept separate.

6.1. RQ1: How often does the primary query have non-positive net utility?

RQ1 concerns the primary low-cost descriptor query under the main cost weighting, $\lambda_T = 1$ and $\lambda_M = 0$. Query utility is decomposed into the observed performance difference, downstream-selector regret, and measured non-FE query cost. Evaluations consumed by query sampling are already absent from the continuation budget of the query branch and are therefore not subtracted again.

Statistical analysis. **TBD** The family-balanced proportion of states with $U_q \leq 0$ is estimated with function-family-aware uncertainty; paired effects decompose performance and cost contributions. State counts are descriptive and do not serve as independent inferential units. Equivalence is evaluated only under the prespecified clustering hierarchy,

resample count, smallest effect of interest, and primary interval or TOST rule.

Answer to RQ1. **TBD** Whether the primary query has net benefit across the evaluated decision states is undetermined. The answer is conditional on the estimated direction, effect size, interval, and observed state coverage, and is restricted to the primary query rather than landscape analysis in general.

6.2. RQ2: Does algorithm-agnostic behavior predict query utility?

The active Decision Model candidates are LDA, Logistic Regression, and Ridge. Their only model-family comparison uses the fixed B3 representation and nested BBOB-train function-family out-of-fold (OOF) mean decision utility. AUROC, Average Precision, and Spearman correlation are auxiliary measures, and continuous-utility RMSE is defined only for Ridge. After this comparison, the selected estimator family is fitted separately to B3 and T0, with a threshold derived from the corresponding complete BBOB-train family-OOF scores. BBOB-validation is evaluation-only and cannot influence preprocessing, model-family choice, input-group comparison, or threshold estimation.

Statistical analysis. **TBD** The paired outer-family BBOB-train OOF comparison for B3, uses family-level effect sizes, confidence intervals, and Holm-corrected candidate contrasts. Conditional on that single model-family selection, the same estimator family defines the B3–T0 comparison and the separate frozen BBOB-validation evaluation. Dependence-aware intervals use the prespecified resampling unit and replication count; auxiliary score measures and BBOB-validation cannot supersede the utility-based BBOB-train selection rule.

Table 6

Decision Model selection and evaluation structure for RQ2. Candidate models are selected using B3 and nested BBOB-train family-OOF utility only. The selected estimator family is subsequently fitted separately to B3 and T0; BBOB-validation is reserved for evaluating these fixed choices without refitting or reselection.

Analysis	Evaluation population	Estimator and representation	Mean utility	AUROC / AP	Spearman / RMSE	Inferential role
Model-family selection	BBOB-train nested family OOF	LDA with B3	–	–	–	Candidate comparison
Model-family selection	BBOB-train nested family OOF	Logistic Regression with B3	–	–	–	Candidate comparison
Model-family selection	BBOB-train nested family OOF	Ridge with B3	–	–	–	Candidate comparison
Representation comparison	BBOB-train family OOF	Selected estimator, B3 versus T0	–	–	–	Paired B3–T0 effect
Frozen evaluation	BBOB-validation	Selected estimator with B3	–	–	–	Evaluation only
Frozen evaluation	BBOB-validation	Selected estimator with T0	–	–	–	Evaluation only

TBD: model-selection, paired representation, and frozen-evaluation estimates with family-level intervals.

RMSE is reported only when the estimator is Ridge; otherwise it is not applicable. Frozen evaluation includes the paired B3–T0 effect.

TBD
(a) Outer-family OOF decision utility for LDA, Logistic Regression, and Ridge with B3; (b) paired B3–T0 effects for the selected estimator family on BBOB-train family-OOF predictions and on the separate frozen BBOB-validation evaluation. Held-out families are the replication units.

Figure 3: B3 model-family comparison and fixed B3–T0 evaluation. The BBOB-validation panel has an evaluation-only role and cannot revise the model choice made from BBOB-train utility.

Answer to RQ2. **TBD** Whether algorithm-agnostic behavior adds decision value beyond FE ratio is undetermined and depends on the paired B3–T0 effects and intervals. BBOB validation is reserved for evaluating only the frozen models; an interval that does not support a stable difference yields an inconclusive answer.

6.3. RQ3: Does Decision-before-Feature reduce unhelpful query calls?

RQ3 retains all eight prespecified comparison roles but reports six distinct outcome rows in the primary population. The native path simultaneously represents Never Query and the training-derived SBS, and the fixed first-opportunity query path simultaneously represents Always Query and Traditional AAS. Random Analysis, the Time-only Controller, and Decision-before-Feature provide the remaining online trajectories; VBS is a per-problem hindsight reference. Slash-labelled pairs are counted once in summaries and inferential comparisons. VBS is neither deployable nor the statewise best observed action.

Statistical analysis. **TBD** The matched function-family Friedman comparison, Holm-adjusted paired Wilcoxon contrasts, effect sizes, and confidence intervals jointly cover mean decision utility as defined in Eq. (37), final error, runtime, and both the rate and accumulated cost of calls with $U_q \leq 0$. The dependence-aware interval construction and equivalence rule are specified independently of the

observed policy differences. Omnibus and follow-up tests use each distinct applicable outcome once; N/A roles are omitted endpoint by endpoint, and neither coincident pair is duplicated.

Answer to RQ3. **TBD** Whether Decision-before-Feature reduces calls with non-positive observed utility is undetermined. A supported answer requires an interval-supported reduction together with the equal-budget effects on final optimization performance and runtime; call rate alone is insufficient.

Table 7

Equal-budget RQ3 comparison. Six distinct outcome rows represent all eight prespecified comparison roles; slash-labelled rows combine definitions that are analytically identical in the primary population and are counted once. Panel A combines optimization performance with resource use. Panel B characterizes query-call quality, including the frequency, within-call rate, and accumulated utility loss of calls whose observed utility is non-positive.

<i>Panel A: optimization performance and resource use</i>						
Outcome row (comparison roles)	Mean decision utility	Final error	Runtime	Total FE	Query FE	Query-call rate
Never Query / SBS	0 / N/A	–	–	–	0 / N/A	0 / N/A
Always Query / Traditional AAS	–	–	–	–	–	–
Random Analysis	–	–	–	–	–	–
VBS	N/A	–	N/A	–	N/A	N/A
Time-only Controller	–	–	–	–	–	–
Decision-before-Feature	–	–	–	–	–	–

<i>Panel B: query-call quality</i>						
Policy	$U_q \leq 0$ calls	Within-call rate	Accumulated utility loss	Beneficial-state capture	Utility capture	Precision
Always Query / Traditional AAS	–	–	–	–	–	–
Random Analysis	–	–	–	–	–	–
Time-only Controller	–	–	–	–	–	–
Decision-before-Feature	–	–	–	–	–	–

TBD: numerical estimates and matched family-level intervals for Panels A and B.
Never Query / SBS and VBS are omitted from Panel B because call-conditional quantities are undefined for static references and Never Query makes no calls.
A slash-labelled row denotes one shared outcome, not an average of two estimates; it enters each applicable analysis once.

TBD

Final optimization loss, measured runtime, and query-call rate for each distinct primary-population outcome under equal FE budgets, with function-family-level uncertainty rather than a frontier based only on point estimates.

Figure 4: Cost–performance comparison among the equal-budget outcomes. Coincident comparison roles are shown once.

6.4. RQ4: Does the frozen controller generalize beyond BBOB training families?

RQ4 reserves the fully frozen procedure for held-out BBOB-validation families, CEC2017, CEC2022, and an engineering-problem evaluation to be specified prospectively. None of these evaluation sets can refit the Selection Reference, Decision Model, preprocessing, representation, or OOF utility threshold. Coverage and query extraction failures are part of the suite-specific evaluation rather than grounds for silently excluding problems.

Statistical analysis. **TBD** Suite-specific problem- or function-family-level paired effects, confidence intervals, and multiplicity-adjusted benchmark contrasts evaluate the frozen policy. Query-extraction failure sensitivity is separate from the primary comparison, and heterogeneous effect directions or coverage preclude a pooled generalization claim.

Answer to RQ4. **TBD** Held-out BBOB performance and cross-benchmark generalization are undetermined. Any eventual answer is suite-specific, distinguishes BBOB validation from CEC2017, CEC2022, and engineering evaluations, and remains conditional on query-extraction coverage.

Table 8

Suite-specific evaluation of the frozen controller for RQ4. Coverage and query-extraction failure rates accompany utility and final optimization performance so that estimates remain conditional on observed evaluation coverage. Paired effects compare Decision-before-Feature with Never Query, Always Query, and the Time-only Controller, in that order.

Evaluation set	Scope	Problem–dimension–seed coverage	Mean utility	Final error	Query failures	Paired effect and interval
BBOB-validation	Held-out BBOB function families	–	–	–	–	–
CEC2017	Cross-benchmark evaluation	–	–	–	–	–
CEC2022	Cross-benchmark evaluation	–	–	–	–	–
Engineering problems	Problem set to be specified prospectively	–	–	–	–	–

TBD: suite-specific coverage, performance, failure, and matched-effect estimates.

TBD

Suite-specific paired effects and confidence intervals for the frozen Decision-before-Feature controller relative to Never Query, Always Query, and the Time-only Controller. Held-out BBOB families are shown separately from CEC and engineering evaluations.

Figure 5: Held-out-family and cross-benchmark evaluation of the frozen controller. Separate panels preserve differences in benchmark scope and data coverage.

6.5. RQ5: Which behavior groups are associated with query utility, and when?

RQ5 reuses the estimator family selected with B3 in RQ2. That estimator family is fitted separately to T0, B1, B2, and B3 using identical samples, function-family splits, train-only preprocessing, and an input-specific BBOB-train family-OOF threshold. The ablation neither repeats model-family selection nor changes the primary B3 representation. T0 contains only FE ratio; B1, B2, and B3 contain the prespecified 19-, 25-, and 31-input permutation-invariant behavior representations. Coefficients, discriminant directions, and associations characterize the fitted rule and are not causal effects.

Statistical analysis. **TBD** The paired outer-family effects and multiplicity-adjusted contrasts compare T0/B1/B2/B3 for the single estimator family selected in RQ2. Interpretation also depends on the sign and magnitude stability of standardized coefficients or discriminant directions across outer folds and families, Spearman associations with observed utility, uncertainty around the fitted maturity–utility curve, and the prespecified dimension and query-configuration analyses. Visual shape alone is insufficient evidence of non-monotonicity, and the ablation cannot repeat model or representation selection.

Answer to RQ5. **TBD** Whether any behavior increment is informative beyond optimization stage is undetermined. Such an interpretation requires agreement among the paired ablation, parameter or discriminant stability, and held-out analyses. Representation- or family-dependent effects instead imply a correspondingly restricted answer; none can replace B3 through post hoc selection.

Table 9

Prespecified input-group ablation for RQ5. The estimator family is the one selected with B3 in RQ2 and is fitted separately to each representation. All groups retain their own BBOB-train family-OOF threshold; the comparison does not select a replacement for B3.

Group	Added information	Inputs	Mean utility	AUROC / AP	Spearman / Ridge RMSE	Paired difference from T0
T0	Optimization stage (FE ratio) only	1	–	–	–	–
B1	Progress, diversity, set and fitness-distribution change, stagnation, convergence, and elite concentration	19	–	–	–	–
B2	B1 plus longitudinal centroid, covariance, fitness-spread, and diversity recovery summaries	25	–	–	–	–
B3	B2 plus set-motion proxies and the three operational maturity variables	31	–	–	–	–

TBD: paired family-level ablation estimates, auxiliary scores, and intervals for all four representations.

The paired difference from T0 concerns decision utility. RMSE is not applicable unless the RQ2-selected estimator is Ridge.

TBD

Fold-stable standardized coefficients for Logistic Regression or Ridge, or standardized discriminant directions for LDA, together with the prespecified association between behavior-derived search maturity and observed query utility. Uncertainty and direction changes are separated by family, dimension, and query configuration.

Figure 6: Behavior interpretation and search-maturity–utility analysis. Displayed coefficients or discriminant directions describe the fitted decision score and do not imply causal effects.

6.6. Prespecified sensitivity analyses and diagnostics

These analyses qualify the RQ1–RQ5 conclusions without defining an additional research question or changing the primary target. Each alternative query configuration traverses the same complete query-specific sequence of action selection, utility estimation, decision modelling, and frozen evaluation as the primary query. The non-primary λ_T values do not replace $\lambda_T = 1$. Action-relation indicators define descriptive strata only, and downstream-selector variants neither enter the Decision Model nor replace the policy baselines.

Table 10

Prespecified sensitivity analyses and downstream-selector diagnostics. Interpretation is restricted to the listed levels and cannot revise the primary cost weighting, input representation, or policy comparison.

Analysis	Prespecified levels	Quantities	Interpretation boundary
Query configuration	Primary low-cost descriptor query, standard <i>pflacco</i> query, and broad <i>pflacco</i> query	Query-specific action selection, utility, decision-model, policy, and frozen-evaluation estimates	Agreement indicates robustness only across these three configurations; disagreement indicates representation or cost dependence
λ_T sensitivity	0, 0.25, 0.5, 1, and 2, with $\lambda_M = 0$	Utility distributions, policy summaries, effect sizes, and intervals	$\lambda_T = 1$ remains the primary conclusion basis
Action-relation strata	Selected action equals the default action; selected action equals the prefix algorithm; population handoff is required	Counts, utility components, selector regret, and policy effects	Descriptive stratification only; no action-relation indicator enters the Decision Model
Downstream-selector inputs	State information only, query descriptors only, and their combination; remaining budget is included in every variant	Cross-fitted action-loss regression performance, selected-action loss, and selector regret	Downstream-selector diagnosis only; these variants are not Decision Model candidates or substitutes for policy baselines

TBD: estimates and intervals for all prespecified sensitivity and diagnostic comparisons.

7. Discussion

The present study separates two decisions that are usually collapsed in feature-based algorithm selection: whether to execute an evaluated fixed landscape query and, conditional on execution, which portfolio action to use. This formulation is methodological and does not by itself establish that the query should be skipped or executed in any empirical regime.

7.1. When is the evaluated query worth executing?

TBD The statewise utility distribution has not yet been estimated. Its eventual interpretation is conditional on the evaluated query, utility definition, effect direction, interval, and state coverage, and cannot be extended from the fixed low-cost query to landscape analysis in general.

7.2. Information in algorithm-agnostic search behavior

TBD The RQ2 comparison between T0 and the behavior groups remains undetermined pending nested BBOB-train family-OOF estimates and frozen evaluation. Decision utility, AUROC, Average Precision, Spearman correlation, and Ridge RMSE support different aspects of the comparison; disagreement among them precludes a single undifferentiated predictive claim.

7.3. End-to-end resource and performance trade-offs

TBD The practical resource–performance trade-off is undetermined. Its interpretation requires joint evidence from equal-total-FE final performance, measured query and inference time, query-call rate, and the cost of calls with non-positive observed utility. Fewer calls alone would not establish better resource-aware optimization.

7.4. Held-out families and cross-benchmark scope

TBD Effects on held-out BBOB, CEC2017, CEC2022, and engineering problems remain unestimated and provide

no current evidence of generalization. Any eventual interpretation is suite-specific; query-extraction failures, divergent effect directions, or incomplete coverage narrow its scope.

7.5. Search maturity and behavior interpretation

TBD The RQ5 ablation, coefficient or discriminant-direction stability, and maturity–utility association remain unestimated. Coefficients, discriminant directions, and correlations have a predictive rather than causal interpretation; a non-monotonic relationship would require support from its prespecified fit, interval, and stability analysis.

7.6. Limitations and validity boundaries

The primary estimand is query-specific: it concerns `descriptor_cheap`, a 16-dimensional custom low-cost descriptor query. The prespecified standard and broad *pflacco* queries can assess configuration robustness but cannot justify claims about untested landscape representations. The Selection Reference is a fixed downstream component rather than a claimed algorithm-selection contribution, and its selector regret must remain visible when interpreting query utility. The main Decision Model is trained offline on a BBOB function-family split; the study does not evaluate online controller training as a primary design.

TBD The empirical limitations remain to be quantified through trajectory and action-loss coverage, query-extraction failures, family-level dependence, population-transfer frequency, and sensitivity to the three prespecified query configurations and utility-cost settings.

8. Reproducibility

8.1. Software and experimental configuration

The study uses offline trajectory collection followed by supervised learning. The experimental specification fixes benchmark membership, dimensions, instances, optimizer seeds, portfolio composition, population size, FE budgets, and the dynamic decision-opportunity schedule. Reporting these elements, together with repetitions, variability, and

feature cost, is particularly important for optimization representations [2].

The main implementation targets Python 3.12. The pflacco queries run in an isolated Python 3.11 environment with pflacco==1.2.2, numpy==1.24.3, pandas==2.0.3, scipy==1.10.1, and scikit-learn==1.2.2. The descriptor extractor receives only the previously evaluated LHS sample. It neither imports a benchmark nor calls the objective function, and a failed pflacco feature group is not replaced by a study-specific substitute.

8.2. Random streams

All scientific random streams are constructed from explicit integers through `numpy.random.SeedSequence`. Optimizer streams combine the configured optimizer seed with integer identifiers for the run and stochastic event. Query sampling additionally uses integer stream, function, instance, dimension, and sample-design codes. Population-transfer initialization receives explicit function, instance, and event codes. The Random Analysis repetitions combine a base seed with problem, dimension, optimizer seed, decision FE, and repetition indices. Experimental randomness therefore does not depend on strings or process-specific hash values.

8.3. Observation indexing and state alignment

An observation is uniquely indexed by

$$k_t = (\text{split}, p, F(p), d, a_t, r, e_t), \quad (39)$$

where p uniquely identifies the problem instance, $F(p)$ is its function family, a_t is the prefix algorithm, r is the optimizer seed, and e_t is the exact integer FE count. This index aligns trajectories, behavior summaries, action outcomes, selection-reference predictions, utility labels, and Decision Model observations. Exact FE, rather than a rounded budget ratio, identifies a shared state; the corresponding elapsed-budget covariate is calculated as e_t/B . Nominal milestone ratios and realized native-window spans describe how the state was sampled but do not replace exact-FE alignment.

8.4. Data construction and evaluation dependencies

Complete native-state trajectories and complete-budget terminal outcomes define the common analysis population. Permutation-invariant behavior and the three fixed descriptor representations are evaluated at matched states, after which the four query-adjusted actions provide the statewise losses. Each query configuration has its own cross-fitted selection reference, utility labels, Decision Model, threshold, baseline comparison, and held-out evaluation. Only trajectories with both a complete native history and a terminal outcome enter the corresponding estimand. Query identity, sample design, FE charge, and descriptor definition jointly identify an analysis condition; in particular, the 5% and 10% acquisition conditions are analyzed separately and are never pooled.

8.5. Validity checks

The validity checks address state continuity, representation invariance, FE accounting, and information leakage. For each portfolio optimizer, a restored complete state is compared with uninterrupted native continuation at matched checkpoints, including optimizer-specific internal dynamics rather than only population values. Independent reordering of population members at every evaluated BBOB state tests that the 31 Decision Model inputs, three diagnostic variables, and native-window summaries are permutation invariant.

Query-level checks enforce the three descriptor definitions, the shared sample for the cheap and standard configurations, four non-duplicated actions, continuous remaining budget, query-specific FE accounting, the statewise action-loss transformation, and the utility and selector decompositions. Feature-group failures remain explicit; undefined individual descriptors are imputed only from BBOB-training observations. Model-level checks fit all preprocessing within the applicable training-family fold, restrict the active candidates to LDA, Logistic Regression, and Ridge, and exclude BBOB validation and external suites from model fitting and threshold estimation. The three action relations in Eq. (35) are checked throughout selection, utility calculation, and policy evaluation. These procedures assess implementation validity; they add no observations to RQ1–RQ5.

8.6. Resource measurement

FE accounting is exact at the state level:

$$FE_{\text{skip,opt}} = B - e_t, \quad (40)$$

$$FE_{q,\text{opt}} = B - e_t - FE_q. \quad (41)$$

Behavior extraction reuses evaluated optimizer states and makes no additional objective-function calls. Query time is the sum of sample-evaluation and feature-computation time. Selection latency includes state preparation, model inference, and action choice. Prediction time is measured separately for the T0 and B3 controllers.

All resource measurements use one fixed hardware and software environment. Processor, operating system, Python environments, thread count, batch size, cache state, and the peak-memory procedure accompany the estimates. Timings are decomposed into behavior extraction, T0 and B3 inference, query sampling, descriptor computation, downstream selection, and total wall time. **TBD**

8.7. Recomputation of estimands

The state-level release retains the quantities needed to recompute every term in Eq. (12) and Eq. (18), including the observed losses, normalization and cost terms, selected action, and three action relations. Model summaries are reconstructed from training-family fold assignments, nested and full-training OOF scores, thresholds, and held-out predictions. Policy summaries are reconstructed from the trigger state, query decision, selected action, transition mode, FE use, timing components, and terminal loss.

Equivalence is interpreted only relative to margins and procedures fixed before policy-outcome inspection. Held-out BBOB and each external suite are analyzed separately with every train-derived component frozen. These provisions define reproducibility and do not themselves constitute evidence of held-out or cross-benchmark transfer.

9. Conclusion

This work formulates analysis selection as a decision that precedes the execution of a prespecified fixed landscape query. The proposed offline workflow links full-state optimizer trajectories, permutation-invariant behavior states, statewise observed action losses, a fixed Selection Reference, and a supervised Decision Model while preserving equal total FE budgets. These design elements define the evaluation; they do not substitute for the empirical evidence required by RQ1–RQ5.

TBD Whether the primary query has non-positive utility in a statistically supported subset of states remains undetermined; the required effect size, interval, and query-specific scope are not yet available.

TBD Whether algorithm-agnostic behavior predicts query utility beyond FE ratio alone remains undetermined pending nested function-family OOF evidence and frozen evaluation.

TBD Effects on calls with non-positive utility, mean decision utility, measured runtime, and equal-budget final optimization performance relative to the prespecified policies remain undetermined.

TBD No conclusion is currently supported for held-out BBOB families or any external suite. BBOB-validation, CEC, and engineering-problem generalization remain unverified pending the corresponding estimates and coverage analyses.

TBD Behavior-group and search-maturity associations remain undetermined and, when estimated, have a descriptive interpretation conditional on the prespecified representations.

The contribution is limited to the prespecified query configurations. Consistent evidence from three prespecified configurations would support configuration robustness, not a universal conclusion about all landscape analysis methods or representations.

References

- [1] van der Blom, K., Vermetten, D., 2026. On the influence of the feature computation budget on per-instance algorithm selection for black-box optimization. doi:10.48550/arXiv.2605.04954, arXiv:2605.04954.
- [2] Cenikj, G., Nikolikj, A., Petelin, G., van Stein, N., Doerr, C., Eftimov, T., 2026. A survey of features used for representing black-box single-objective continuous optimization. *Swarm and Evolutionary Computation* 101, 102288. doi:10.1016/j.swevo.2026.102288.
- [3] Cenikj, G., Petelin, G., Doerr, C., Korošec, P., Eftimov, T., 2023. DynamoRep: Trajectory-based population dynamics for classification of black-box optimization problems, in: *Proceedings of the Genetic and Evolutionary Computation Conference*, pp. 813–821. doi:10.1145/3583131.3590401.
- [4] Cenikj, G., Petelin, G., Seiler, M., Cenikj, N., Eftimov, T., 2025. Landscape features in single-objective continuous optimization: Have we hit a wall in algorithm selection generalization? *Swarm and Evolutionary Computation* 94, 101894. doi:10.1016/j.swevo.2025.101894.
- [5] Eimer, T., Biedenkapp, A., Reimer, M., Adriansen, S., Hutter, F., Lindauer, M., 2021. DACBench: A benchmark library for dynamic algorithm configuration, in: *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence, International Joint Conferences on Artificial Intelligence Organization*, pp. 1668–1674. doi:10.24963/ijcai.2021/230.
- [6] Guo, H., Ma, Y., Ma, Z., Chen, J., Zhang, X., Cao, Z., Zhang, J., Gong, Y.J., 2024. Deep reinforcement learning for dynamic algorithm selection: A proof-of-principle study on differential evolution. *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 54, 4247–4259. doi:10.1109/TSMC.2024.3374889.
- [7] Guo, Q., Wang, H., Tian, Y., 2025. Automated algorithm selection for black-box optimization using light gradient boosting machine. *Swarm and Evolutionary Computation* 98, 102071. doi:10.1016/j.swevo.2025.102071.
- [8] Hayward, L., Engelbrecht, A., 2025. Determining metaheuristic similarity using behavioral analysis. *IEEE Transactions on Evolutionary Computation* 29, 262–274. doi:10.1109/TEVC.2023.3346672.
- [9] Hayward, L., Engelbrecht, A.P., 2026. Survey and analysis of metaheuristic search behavior characterization: A case study on particle swarm optimization variants. *Swarm Intelligence* 20, 1–33. doi:10.1007/s11721-025-00254-1.
- [10] Janković, A., Doerr, C., 2019. Adaptive landscape analysis, in: *Proceedings of the Genetic and Evolutionary Computation Conference Companion, ACM*, pp. 2032–2035. doi:10.1145/3319619.3326905.
- [11] Jankovic, A., Vermetten, D., Kostovska, A., de Nobel, J., Eftimov, T., Doerr, C., 2022. Trajectory-based algorithm selection with warm-starting, in: *2022 IEEE Congress on Evolutionary Computation (CEC)*, IEEE, pp. 1–8. doi:10.1109/CEC55065.2022.9870222.
- [12] Kerschke, P., Trautmann, H., 2019. Automated algorithm selection on continuous black-box problems by combining exploratory landscape analysis and machine learning. *Evolutionary Computation* 27, 99–127. doi:10.1162/evco_a_00236.
- [13] Korošec, P., Eftimov, T., 2024. Opt2Vec—a continuous optimization problem representation based on the algorithm’s behavior: A case study on problem classification. *Information Sciences* 680, 121134. doi:10.1016/j.ins.2024.121134.
- [14] Kostovska, A., Jankovic, A., Vermetten, D., de Nobel, J., Wang, H., Eftimov, T., Doerr, C., 2022. Per-run algorithm selection with warm-starting using trajectory-based features, in: *Parallel Problem Solving from Nature – PPSN XVII*, Springer, pp. 46–60. doi:10.1007/978-3-031-14714-2_4.
- [15] Gomes Pereira de Lacerda, M., de Araujo Pessoa, L.F., Buarque de Lima Neto, F., Ludermir, T.B., Kuchen, H., 2021. A systematic literature review on general parameter control for evolutionary and swarm-based algorithms. *Swarm and Evolutionary Computation* 60, 100777. doi:10.1016/j.swevo.2020.100777.
- [16] Malan, K.M., Engelbrecht, A.P., 2013. A survey of techniques for characterising fitness landscapes and some possible ways forward. *Information Sciences* 241, 148–163. doi:10.1016/j.ins.2013.04.015.
- [17] Mbasso, W.F., Harrison, A., Liu, Z., Mamadiyarov, Z., Al-Gahtani, S.F., Elbarbary, Z.M.S., 2026. How do metaheuristics exploit? a particle-level behavioral study across the CEC 2025 benchmark functions. *Evolutionary Intelligence* 19, 77. doi:10.1007/s12065-026-01190-7.
- [18] Mersmann, O., Bischl, B., Trautmann, H., Preuss, M., Weihs, C., Rudolph, G., 2011. Exploratory landscape analysis, in: *Proceedings of the 13th Annual Conference on Genetic and Evolutionary Computation, ACM*, pp. 829–836. doi:10.1145/2001576.2001690.
- [19] Ochoa, G., Malan, K.M., Blum, C., 2021. Search trajectory networks: A tool for analysing and visualising the behaviour of metaheuristics. *Applied Soft Computing* 109, 107492. doi:10.1016/j.asoc.2021.107492.

- [20] Pei, J., Mei, Y., Liu, J., Zhang, M., Yao, X., 2025. Adaptive operator selection for meta-heuristics: A survey. *IEEE Transactions on Artificial Intelligence* 6, 1991–2012. doi:10.1109/TAI.2025.3545792.
- [21] Petelin, G., Cenikj, G., 2025. The pitfalls of benchmarking in algorithm selection: What we are getting wrong, in: *Proceedings of the Genetic and Evolutionary Computation Conference*, ACM. pp. 1181–1189. doi:10.1145/3712256.3726336.
- [22] Prager, R.P., Trautmann, H., 2024. Pflacco: Feature-based landscape analysis of continuous and constrained optimization problems in python. *Evolutionary Computation* 32, 211–216. doi:10.1162/evco_a_00341.
- [23] Renau, Q., Doerr, C., Dreio, J., Doerr, B., 2020. Exploratory landscape analysis is strongly sensitive to the sampling strategy, in: *Parallel Problem Solving from Nature – PPSN XVI*. Springer International Publishing. volume 12270, pp. 139–153. doi:10.1007/978-3-030-58115-2_10.
- [24] Renau, Q., Hart, E., 2024. On the utility of probing trajectories for algorithm-selection, in: *Applications of Evolutionary Computation*, Springer. pp. 98–114. doi:10.1007/978-3-031-56852-7_7.
- [25] Renau, Q., Hart, E., 2025. Algorithm selection with probing trajectories: Benchmarking the choice of classifier model, in: *Applications of Evolutionary Computation*, Springer. pp. 452–468. doi:10.1007/978-3-031-90062-4_28.
- [26] Sallam, K.M., Elsayed, S.M., Sarker, R.A., Essam, D.L., 2017. Landscape-based adaptive operator selection mechanism for differential evolution. *Information Sciences* 418–419, 383–404. doi:10.1016/j.ins.2017.08.028.
- [27] Yang, X., Wang, R., Li, K., Ishibuchi, H., 2025. Meta-black-box optimization for evolutionary algorithms: Review and perspective. *Swarm and Evolutionary Computation* 93, 101838. doi:10.1016/j.swevo.2024.101838.
- [28] Zhou, X., Zhang, X., Gao, W., Wang, H., Ma, Y., 2024. Adaptive multi-population artificial bee colony algorithm based on fitness landscape analysis. *Applied Soft Computing* 164, 111952. doi:10.1016/j.asoc.2024.111952.